

SFWRENG 3SH3: Operating Systems

Lab 3: POSIX Threads

Dr. Borzoo Bonarkdarpour

1 Overall Objective

You are to write a C or C++ concurrent *Shearsort* program using POSIX Threads (PThreads) library. You are NOT to use `fork()` calls, pipes, or message queues in this exercise.

The *Shearsort* is a simple mesh-sorting algorithm that consists of nothing more than alternately sorting rows and columns of the mesh. In particular, it sorts all the **rows** in Phases 1, 3, \dots , $\log_2 \sqrt{N} + 1$, and all the **columns** in Phases 2, 4, \dots , $\log_2 \sqrt{N}$, where N is the total number of elements. The **columns** are sorted so that smaller numbers move upward. The **odd rows** (1, 3, \dots , $\sqrt{N}-1$) are sorted so that smaller numbers move leftward, and the **even rows** (2, 4, \dots , \sqrt{N}) are sorted in reverse order (i.e., so that smaller numbers move rightward). The numbers will appear in a snakelike order after $2 \log_2 \sqrt{N} + 1 = \log_2 N + 1$ phases. An example is given in Figure 1. (A proof that the Shearsort algorithm works can be found in *Introduction to Parallel Algorithms and Architectures: Arrays, Trees, Hypercubes* by F. Thomson Leighton. However, the information about Shearsort in this handout should be sufficient for you to complete the exercise.)

2 Assignment

In this lab assignment you will use shared memory to store ONE copy of the whole 2-dimensional array. Also, you will be required to create a limited number of threads for the sort. For a square array of $n \times n$ elements, ONLY n THREADS are to be created. Each thread is to alternate between row (odd) and column (even) phases. To synchronize the phases properly, you may use either *semaphores* (at most ONE SEMAPHORE PER THREAD), or a more elegant data structure for this purpose, called *condition variables*.

Your program should implement the Shearsort algorithm to handle 16 integers as follows.

1. Read the integers into a 2-dimensional $n \times n$ array in global memory from the file “input.txt”. This can be done by the main thread before creating the n sorting threads.
2. Print the integers in the order entered to `stdout`.
3. Create and initialize the semaphores or condition variable necessary for the algorithm.

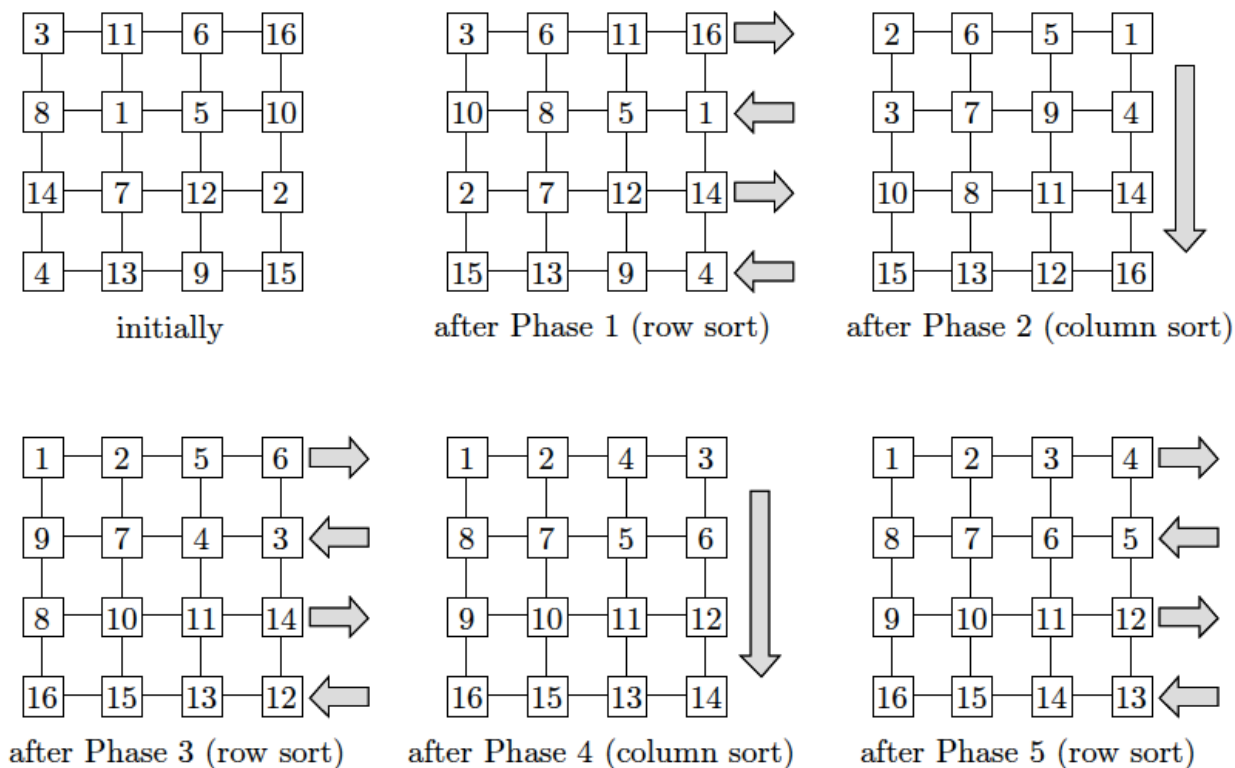


Figure 1: Alternately sorting the rows and columns in Shearsort. The numbers to be sorted appear in a snakelike order after $\log_2 N + 1$ phases. Notice that even rows are always sorted in reverse order. The arrow direction indicates the sorting order (from small to large).

4. Create the n threads to sort the array using Shearsort.
5. Wait for the threads to finish a phase.
6. Print the array of sorted integers to `stdout` and go to the next phase..

Each thread would do the following in each phase of the algorithm.

1. By performing the appropriate number of wait operations on semaphores (or by a proper wait on a condition variable), block until the prior phase (if any) is finished.
2. Sort the row/column in the appropriate order using Bubble Sort.
3. Perform appropriate signal operations to signal the other threads to begin the next phase.

WARNING: You are NOT to use the `sleep()` call or any code (such as a loop counting from 1 to 10000) to introduce any delay in your program. If any such delaying code is found in your program, YOU WILL LOSE %30 OF TOTAL POINTS.

3 POSIX pthreads

POSIX Threads, usually referred to as Pthreads, is an execution model that exists independently from a language, as well as a parallel execution model. It allows a program to control multiple different flows of work that overlap in time. Each flow of work is referred to as a thread, and creation and control over these flows is achieved by making calls to the POSIX Threads API. You can find plenty of online resources to learn about the API. Essentially you need to know how to (1) create threads, (2) terminate threads, and (3) use mutex variables. Here's one resource: <https://computing.llnl.gov/tutorials/pthreads/>

4 Guidelines

You will

- work on this assignment individually
- implement this assignment using C or C++ (you are free to separate your program into multiple source files as you see fit)
- present your implementation and output to your lab TA.

You may present your program in a different section in the same week **only once** throughout the term.